

Homework Assignment #1
Due: Monday, 10 February 2020 at 19h00 (7pm)

Red team, blue team: Identifying political persuasion on Reddit

TAs: Julia Watson (j.watson@mail.utoronto.ca) and Krishnapriya Vishnubhotla (vkpriya@cs)

Introduction

This assignment will give you experience with a social media corpus (i.e., a collection of posts from Reddit), Python programming, part-of-speech (PoS) tags, sentiment analysis, and machine learning with scikit-learn.

Your task is to split posts into sentences, tag them with a PoS tagger that we will provide, gather some feature information from each post, learn models, and use these to classify political persuasion. Sentiment analysis is an important topic in computational linguistics in which we quantify subjective aspects of language. These aspects can range from biases in social media for marketing, to a spectrum of cognitive behaviours for disease diagnosis.

Please check [the course bulletin board](#) for announcements and discussion pertaining to this assignment.

Reddit Corpus

We have curated data from [Reddit](#) by scraping subreddits, using [Pushshift](#), by perceived political affiliation. Table 1 shows the subreddits assigned to each of four categories: left-leaning, right-leaning, center/neutral, and ‘alternative facts’. Although the first three (at least) are often viewed as ordinal segments on a unidimensional spectrum, here we treat these categories as nominal classes. Here, we use the terms ‘post’ and ‘comment’ interchangeably to mean ‘datum’, i.e., a short segment of user-produced text.

Left (598,944)	Center (599,872)	Right (600,002)	Alt (200,272)
twoXChromosomes (7,720,661)	news (2,782,991)	theNewRight (19,466)	conspiracy (6,767,099)
occupyWallStreet (397,538)	politics (60,354,767)	whiteRights (118,008)	911truth (79,868)
lateStageCapitalism (634,962)	energy (416,926)	Libertarian (3,886,156)	
progressive (246,435)	canada (7,225,005)	AskTrumpSupporters (1,007,590)	
socialism (1,082,305)	worldnews (38,851,904)	The_Donald (21,792,999)	
demsocialist (5269)	law (464,236)	new_right (25,166)	
Liberal (151,350)		Conservative (1,929,977)	
		tea_party (1976)	

Table 1: Subreddits assigned to each category, with the *total* posts in each. Since there are over 181M posts, we sample randomly within each category – the resulting number of *available* posts for each category in this assignment is shown on the top row.

These data are stored on the teach.cs servers under `/u/cs401/A1/data/`. To save space, these files should *only* be accessed from that directory (and not copied). All data are in the [JSON](#) format.

Each datum has several fields of interest, including:

- ups:** the integer number of upvotes.
- downs:** the integer number of downvotes.
- score:** $[ups - downs]$
- controversiality:** a combination of the popularity of a post and the ratio between ups and downs.
- subreddit:** the subreddit from which the post was sampled.
- author:** the author ID.
- body:** the main textual message of the post, and our primary interest.
- id:** the unique identifier of the comment.

Your tasks

1 Pre-processing, tokenizing, and tagging [25 marks]

The comments, as given, are not in a form amenable to feature extraction for classification – there is too much ‘noise’. Therefore, the first step is to complete a Python program named `a1_preproc.py`, in accordance with Section 5, that will read subsets of JSON files, and for each comment perform the following steps, in order, on the ‘body’ field of each selected comment:

1. Replace all newline characters with spaces.
2. Replace HTML character codes (i.e., `–`) with their ASCII equivalent (see <http://www.asciitable.com>).
3. Remove all URLs (i.e., tokens beginning with `http` or `www`).
4. Remove duplicate spaces between tokens.
 - Each token must now be separated by a single space.
5. Apply the following steps using spaCy (see below):
 - Tagging: Tag each token with its part-of-speech. A tagged token consists of a word, the ‘/’ symbol, and the tag (e.g., `dog/NN`). See below for information on how to use the tagging module. The tagger can make mistakes.
 - Lemmatization: Replace the token itself with the `token.lemma_`. E.g., `words/NNS` becomes `word/NNS`. If the lemma begins with a dash (‘-’) when the token doesn’t (e.g., `-PRON-` for *I*, just keep the token.
 - Sentence segmentation: Add a newline between each sentence. For this assignment, we will use spaCy’s `sentencizer` component to segment sentences in a post. Remember to also mark the end of the post with a newline (watch out for duplicates!).

spaCy: spaCy is a Python library for natural language processing tasks, especially in information extraction. Here, we **only** use its ability to obtain part-of-speech tags and lemma, along with sentence segmentation. For example:

```
import spacy

nlp = spacy.load('en_core_web_sm', disable=['parser', 'ner'])
sentencizer = nlp.create_pipe("sentencizer")
nlp.add_pipe(sentencizer)

utt = nlp(u"I know words. I have the best words")

for sent in utt.sents:
...     print(sent.text)
...     for token in sent:
...         print(token.text, token.lemma_, token.pos_, token.tag_, token.dep_,
...               token.shape_, token.is_alpha, token.is_stop)
```

Functionality: The `a1_preproc.py` program reads a subset of the (static) input JSON files, retains the fields you care about, including `'id'`, which you'll soon use as a key to obtain pre-computed features, and `'body'`, which is text that you preprocess and replace before saving the result to an output file. To each comment, also add a `cat` field, with the name of the file from which the comment was retrieved (e.g., `'Left'`, `'Alt'`,...).

The program takes three arguments: your student ID (mandatory), the output file (mandatory), and the maximum number of lines to sample from each category file (optional; default=10,000). For example, if you are student 999123456 and want to create `preproc.json`, you'd run:

```
python a1_preproc.py 999123456 -o preproc.json
```

The output of `a1_preproc.py` will be used in Task 2.

Your task: Copy the template from `/u/cs401/A1/code/a1_preproc.py`. There are two functions you need to modify:

1. In `preproc1`, fill out each `if` statement with the associated preprocessing step above.
2. In `main`, replace the lines marked with `TODO` with the code they describe. By default, `args.a1_dir` points to `/u/cs401/A1/`, so load the data from `args.a1_dir`'s data subdirectory.

For this section, you may only use standard Python libraries, *except* for Step 5. For debugging, you are advised to either use a different input folder with your own JSON data, or pass strings directly to `preproc1`.

Subsampling: By default, you should only sample 10,000 lines from each of the Left, Center, Right, and Alt files, for a total of 40,000 lines. From each file, start sampling lines at index $[ID \% \text{len}(X)]$, where ID is your student ID, $\%$ is the modulo arithmetic operator, and $\text{len}(X)$ is the number of comments in the given input file (i.e., `len(data)`, once the JSON parse is done). Use circular list indexing if your start index is too close to the 'end'.

2 Feature extraction [20 marks]

The second step is to complete a Python program named `a1_extractFeatures.py`, in accordance with Section 5, that takes the preprocessed comments from Task 1, extracts features that are relevant to bias detection, and builds an `npz` datafile that will be used to train models and classify comments in Task 3.

For each comment, you need to extract 173 features and write these, along with the category, to a single `NumPy array`. These features are listed below. Several of these features involve counting tokens based on their tags. For example, counting the number of *adverbs* in a comment involves counting the number of tokens that have been tagged as RB, RBR, or RBS. Table 4 explicitly defines some of these features (many of which we have provided as constants in the template); other definitions are available on CDF in `/u/cs401/Wordlists/`. You may copy and modify these files, but do not change their filenames.

1. Number of words in uppercase (≥ 3 letters long)
2. Number of first-person pronouns
3. Number of second-person pronouns
4. Number of third-person pronouns
5. Number of coordinating conjunctions
6. Number of past-tense verbs
7. Number of future-tense verbs
8. Number of commas
9. Number of multi-character punctuation tokens
10. Number of common nouns
11. Number of proper nouns
12. Number of adverbs
13. Number of *wh*- words
14. Number of slang acronyms
15. Average length of sentences, in tokens
16. Average length of tokens, excluding punctuation-only tokens, in characters
17. Number of sentences.
18. Average of AoA (100-700) from Bristol, Gilhooly, and Logie norms
19. Average of IMG from Bristol, Gilhooly, and Logie norms
20. Average of FAM from Bristol, Gilhooly, and Logie norms
21. Standard deviation of AoA (100-700) from Bristol, Gilhooly, and Logie norms
22. Standard deviation of IMG from Bristol, Gilhooly, and Logie norms
23. Standard deviation of FAM from Bristol, Gilhooly, and Logie norms
24. Average of V.Mean.Sum from Warringer norms
25. Average of A.Mean.Sum from Warringer norms
26. Average of D.Mean.Sum from Warringer norms
27. Standard deviation of V.Mean.Sum from Warringer norms
28. Standard deviation of A.Mean.Sum from Warringer norms
29. Standard deviation of D.Mean.Sum from Warringer norms
- 30-173. LIWC/Receptiviti features

Note: All of the provided wordlists contain tokens in lowercase. After extracting feature 1 above, you may convert the tokens in the comment to lowercase as well. Take care to modify only the text and not the PoS tags, i.e., *Dog/NN* should become *dog/NN* and not *dog/nn*.

Functionality: The `a1_extractFeatures.py` program reads a preprocessed JSON file and extracts features for each comment therein, producing and saving a $D \times 174$ NumPy array, where the i^{th} row is the features for the i^{th} comment, followed by an integer for the class (0: Left, 1: Center, 2: Right, 3: Alt), as per the `cat` JSON field.

The program takes two arguments: the input filename (i.e., the output of `a1_preproc`), and the output filename. For example, given input `preproc.json` and the desired output `feats.npz`, you'd run:

```
python a1_extractFeatures.py -i preproc.json -o feats.npz
```

The output of `a1_extractFeatures.py` will be used in Task 3.

Your task: Copy the template from `/u/cs401/A1/code/a1_extractFeatures.py`. There are two functions you need to modify:

1. In `extract1`, extract each the first 29 of the aforementioned features from the input string. Features 30-173 should be extracted in `extract2`.
2. In `main`, call `extract1` and `extract2` on each datum and add the results (+ the class) to the `feats` array.

When your feature extractor works to your satisfaction, build `feats.npz`, from all input data.

Norms: Lexical norms are aggregate subjective scores given to words by a large group of individuals. Each type of norm assigns a numerical value to each word. Here, we use two sets of norms:

Bristol+GilhoolyLogie: These are found in `/u/cs401/Wordlists/BristolNorms+GilhoolyLogie.csv`, specifically the fourth, fifth, and sixth columns. These measure the Age-of-acquisition (AoA), imageability (IMG), and familiarity (FAM) of each word, which we can use to measure lexical complexity. More information can be found, for example, [here](#).

Warringer: These are found in `/u/cs401/Wordlists/Ratings_Warriner_et_al.csv`, specifically the third, sixth, and ninth columns. These norms measure the valence (V), arousal (A), and dominance (D) of each word, according to the [VAD](#) model of human affect and emotion. More information on this particular data set can be found [here](#).

When you compute features 18-29, only consider those words that exist in the respective norms file.

Assume the default value for all of the above features to be zero. Treat the mean and standard deviation of zero words to be zero.

LIWC/Receptiviti: The Linguistic Inquiry & Word Count (LIWC) tool has been a standard in a variety of NLP research, especially around authorship and sentiment analysis. This tool provides 85 measures mostly related to word choice; more information can be found [here](#). The company [Receptiviti](#) provides a superset of these features, which also includes 59 measures of personality derived from text. The company has graciously donated access to its API for the purposes of this course.

To simplify things, we have already extracted these 144 features for you. Simply copy the pre-computed features from the appropriate *uncompressed* `npz` files stored in `/u/cs401/A1/feats/`. Specifically:

1. Comment IDs are stored in `_IDs.txt` files (e.g., `Alt_Ids.txt`). When processing a comment, find the index (row) i of the ID in the appropriate ID text file, for the category, and copy the 144 elements of that row from the associated `_feats.dat.npz` file.
2. The file `feats.txt` provides the names of these features, in the order provided. For this assignment, these names will suffice as to their meaning, but you are welcome to obtain your own API license from [Receptiviti](#) in order to get access to their documentation.

3 Experiments and classification [30 marks]

The third step is to use the features extracted in Task 2 to classify comments using the scikit-learn machine learning package. Here, you will modify various hyper-parameters and interpret the results analytically. As everyone has different slices of the data, there are no expectations on overall *accuracy*, but you are expected to discuss your findings with scientific rigour. Copy the template from `/u/cs401/A1/code/a1_classify.py` and complete the `main` body and the functions for the following experiments according to the specifications therein.

The program takes two arguments: the input feature file (the output of `a1_extractFeatures`), and an output directory.

```
python a1_classify.py -i feats.npz -o .
```

You should create the output directory if it doesn't already exist. In `main`, you are expected to load the data from the input file, partition the input into a train and test set, and call the experiment functions in order. Use the `train_test_split` method to split the data into a random 80% for training and 20% for testing. For part 3.3, use the entire loaded data set.

3.1 Classifiers

Train the following 5 classifiers (see hyperlinks for API) with `fit(X_train, y_train)`:

1. SGDClassifier: support vector machine with a linear kernel.
2. GaussianNB: a Gaussian naive Bayes classifier.
3. RandomForestClassifier: with a maximum depth of 5, and 10 estimators.
4. MLPClassifier: A feed-forward neural network, with $\alpha = 0.05$.
5. AdaBoostClassifier: with the default hyper-parameters.

Here, `X_train` is the first 173 columns of your training data, and `y_train` is the last column. Obtain predicted labels with these classifiers using `predict(X_test)`, where `X_test` is the first 173 columns of your testing data. Obtain the 4×4 confusion matrix C using `confusion_matrix`. Given that the element at row i , column j in C (i.e., $c_{i,j}$) is the number of instances belonging to class i that were classified as class j , compute the following manually, using the associated function templates:

Accuracy : the total number of correctly classified instances over all classifications: $A = \frac{\sum_i c_{i,i}}{\sum_{i,j} c_{i,j}}$.

Recall : for each class κ , the fraction of cases that are truly class κ that were classified as κ , $R(\kappa) = \frac{c_{\kappa,\kappa}}{\sum_j c_{\kappa,j}}$.

Precision : for each class κ , the fraction of cases classified as κ that truly are κ , $P(\kappa) = \frac{c_{\kappa,\kappa}}{\sum_i c_{i,\kappa}}$.

Write the results to the text file `a1.3.1.txt` in the output directory. You must write to file using the format strings provided in the template. **If you do not follow the format, you may receive a mark of zero.** For each classifier, you will print the accuracy, recall, precision, and confusion matrix. You may include a written analysis if you are so inclined, but only after the results.

3.2 Amount of training data

Many researchers attribute the success of modern machine learning to the sheer volume of data that is now available. Modify the amount of data that is used to train your preferred classifier from above in five increments: 1K, 5K, 10K, 15K, and 20K. These can be sampled arbitrarily from the training set in Section 3.1. Using only the classification algorithm with the highest accuracy from Section 3.1, report the accuracies of the classifier to the file `a1.3.2.txt` using the format string provided in the template. On one or more lines following the reported accuracies, comment on the changes to accuracy as the number of training samples increases, including at least two sentences on a possible explanation. Is there an expected trend? Do you see such a trend? Hypothesize as to why or why not.

3.3 Feature analysis

Certain features may be more or less useful for classification, and too many can lead to overfitting or other problems. Here, you will select the best features for classification using `SelectKBest` according to the `f_classif` metric as in:

```
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif

selector = SelectKBest(f_classif, you_figure_it_out)
X_new = selector.fit_transform(X_train, y_train)
pp = selector.pvalues_
```

In the example above, `pp` stores the p -value associated with doing a χ^2 statistical test on each feature. A smaller value means the associated feature better separates the classes. Do this:

1. For the 32k training set and each number of features $k = \{5, 50\}$, find the best k features according to this approach. Write the associated p -values to `a1.3.3.txt` using the format strings provided.
2. Train the best classifier from section 3.1 for each of the 1K training set and the 32K training set, using only the best $k = 5$ features. Write the accuracies on the full test set of both classifiers to `a1.3.3.txt` using the format strings provided.
3. Extract the indices of the top $k = 5$ features using the 1K training set and take the intersection with the $k = 5$ features using the 32K training set. Write using the format strings provided.
4. Format the top $k = 5$ feature indices extracted from the 32K training set to file using the format string provided.
5. Following the above, answer the following questions:
 - (a) Provide names for the features found in the above intersection of the top $k = 5$ features. If any, provide a possible explanation as to why these features may be especially useful.
 - (b) Are p -values generally higher or lower given more or less data? Why or why not?
 - (c) Name the top 5 features chosen for the 32K training case. Hypothesize as to why those particular features might differentiate the classes.

3.4 Cross-validation

Many papers in machine learning stick with a single subset of data for training and another for testing (occasionally with a third for validation). This may not be the most honest approach. Is the best classifier from Section 3.1 *really* the best? For each of the classifiers in Section 3.1, run 5-fold cross-validation given all the initially available data. Specifically, use `KFold`. **Set the shuffle argument to true.**

For each fold, obtain accuracy on the test partition after training on the rest for each classifier. Report the mean accuracy of each classifier across the 5 folds in the order specified in 3.1 to `a1.3.4.txt` using the format strings provided. Next, determine whether the accuracy of your best classifier, across the 5 folds, is *significantly* better than any others. I.e., given vectors `a` and `b`, one for each classifier, containing the accuracy values for each of the respective 5 folds, obtain the p -value from the output `S`, below:

```
from scipy import stats
S = stats.ttest_rel(a, b)
print(S.pvalue)
```

You should have 4 p -values. Report them using the provided format string in the same order as the accuracies, excluding the self-comparison. For example, if the best classifier from 3.1 was the `RandomForestClassifier`, then the p -values should be reported in the order: 1 vs. 3, 2 vs. 3, 4 vs. 3, 5 vs. 3.

4 Bonus [15 marks]

We will give up to 15 bonus marks for innovative work going substantially beyond the minimal requirements. These marks can make up for marks lost in other sections of the assignment, but your overall mark for this assignment cannot exceed 100%. The obtainable bonus marks will depend on the complexity of the undertaking, and are at the discretion of the marker. Importantly, your bonus work should not affect our ability to mark the main body of the assignment in any way.

You may decide to pursue any number of tasks of your own design related to this assignment, although you should consult with the instructor or the TA before embarking on such exploration. Certainly, the rest of the assignment takes higher priority. Some ideas:

- **Identify words that the PoS tagger tags incorrectly** and add code that fixes those mistakes. Does this code introduce new errors elsewhere? E.g., if you always tag *dog* as a noun to correct a mistake, you will encounter errors when *dog* should be a verb. How can you mitigate such errors?
- Explore **alternative features** to those extracted in Task 2. What other kinds of variables would be useful in distinguishing affect? Consider, for example, the Stanford Deep Learning for Sentiment Analysis. **Test your features empirically as you did in Task 3 and discuss your findings.**
- Explore **alternative classification methods** to those used in Task 3. Explore **different hyper-parameters**. Which hyper-parameters give the best empirical performance, and why?
- Learn about **topic modelling** as in [latent Dirichlet allocation](#). Are there topics that have an effect on the accuracy of the system? E.g., is it easier to tell how someone feels about politicians or about events? People or companies? As there may be class imbalances in the groups, how would you go about evaluating this? Go about evaluating this.

5 General specifications

As part of grading your assignment, the grader may run your programs and/or python files on test data and configurations that you have not previously seen. This may be partially done automatically by scripts. It is therefore important that each of your programs precisely meets all the specifications, including its name and the names of the files and functions that it uses. **A program that cannot be evaluated because it varies from specifications will receive zero marks on the relevant sections.**

The flag `--a1_dir` can be used to specify an alternate location than `/u/cs401/A1` to load data from, though the submitted files should be generated on the CDF machines. Do **not** hardwire the absolute address of your home directory within the program; the grader does not have access to this directory.

All your programs must contain adequate internal documentation to be clear to the graders.

We use Python version 3.7.

6 Submission requirements

This assignment is submitted electronically. You should submit:

1. All your code for `a1_preproc.py`, `a1_extractFeatures.py`, and `a1_classify.py` (including helper files, if any).
2. `a1_3.1.txt`: Report on classifiers.
3. `a1_3.2.txt`: Report on the amount of training data.
4. `a1_3.3.txt`: Report on feature analysis.
5. `a1_3.4.txt`: Report on 5-fold cross-validation.
6. Any lists of words that you modified from the original version.

In another file called *ID* (use the template on the course web page), provide the following information:

1. your first and last name.
2. your student number.
3. your CDF/teach.cs login id.
4. your preferred contact email address.
5. whether you are an undergraduate or graduate.
6. this statement: *By submitting this file, I declare that my electronic submission is my own work, and is in accordance with the University of Toronto Code of Behaviour on Academic Matters and the Code of Student Conduct, as well as the collaboration policies of this course.*

You do not need to hand in any files other than those specified above. Submit your assignment on MarkUs. Do **not** `tar` or `compress` your files, and do not place your files in subdirectories. Do not format your discussion as a PDF or Word document — use plain text only.

7 Working outside the lab

If you want to do some or all of this assignment on your laptop or home computer, for example, you will have to do the extra work of downloading and installing the requisite software and data. If you take this route, you take on all associated risks. You are strongly advised to upload regular backups of your work to CDF/teach.cs, so that if your home machine fails or proves to be inadequate, you can immediately continue working on the assignment at CDF/teach.cs. When you have completed the assignment, you should try your programs out on CDF/teach.cs to make sure that they run correctly there. **Any component that does not work on CDF will get zero marks.**

Appendix: Tables

Tag	Name	Example
CC	Coordinating conjunction	<i>and</i>
CD	Cardinal number	<i>three</i>
DT	Determiner	<i>the</i>
EX	Existential <i>there</i>	<i>there [is]</i>
FW	Foreign word	<i>d'oeuvre</i>
IN	Preposition or subordinating conjunction	<i>in, of, like</i>
JJ	Adjective	<i>green, good</i>
JJR	Adjective, comparative	<i>greener, better</i>
JJS	Adjective, superlative	<i>greenest, best</i>
LS	List item marker	<i>(1)</i>
MD	Modal	<i>could, will</i>
NN	Noun, singular or mass	<i>table</i>
NNS	Noun, plural	<i>tables</i>
NNP	Proper noun, singular	<i>John</i>
NNPS	Proper noun, plural	<i>Vikings</i>
PDT	Predeterminer	<i>both [the boys]</i>
POS	Possessive ending	<i>'s, '</i>
PRP	Personal pronoun	<i>I, he, it</i>
PRP\$	Possessive pronoun	<i>my, his, its</i>
RB	Adverb	<i>however, usually, naturally, here, good</i>
RBR	Adverb, comparative	<i>better</i>
RBS	Adverb, superlative	<i>best</i>
RP	Particle	<i>[give] up</i>
SYM	Symbol (mathematical or scientific)	<i>+</i>
TO	<i>to</i>	<i>to [go] to [him]</i>
UH	Interjection	<i>uh-huh</i>
VB	Verb, base form	<i>take</i>
VBD	Verb, past tense	<i>took</i>
VBG	Verb, gerund or present participle	<i>taking</i>
VBN	Verb, past participle	<i>taken</i>
VBP	Verb, non-3rd-person singular present	<i>take</i>
VBZ	Verb, 3rd-person singular present	<i>takes</i>
WDT	<i>wh</i> -determiner	<i>which</i>
WP	<i>wh</i> -pronoun	<i>who, what</i>
WP\$	Possessive <i>wh</i> -pronoun	<i>whose</i>
WRB	<i>wh</i> -adverb	<i>where, when</i>

Table 2: The Penn part-of-speech tagset—words

Tag	Name	Example
#	Pound sign	£
\$	Dollar sign	\$
.	Sentence-final punctuation	!, ?, .
,	Comma	
:	Colon, semi-colon, ellipsis	
(Left bracket character	
)	Right bracket character	
"	Straight double quote	
'	Left open single quote	
“	Left open double quote	
'	Right close single quote	
”	Right close double quote	

Table 3: The Penn part-of-speech tagset—punctuation

First person:

I, me, my, mine, we, us, our, ours

Second person:

you, your, yours, u, ur, urs

Third person:

he, him, his, she, her, hers, it, its, they, them, their, theirs

Future Tense:

'll, will, gonna, going+to+VB

Common Nouns:

NN, NNS

Proper Nouns:

NNP, NNPS

Adverbs:

RB, RBR, RBS

wh-words :

WDT, WP, WP\$, WRB

Modern slang acronyms:

smh, fwb, lmfaol, lmao, lms, tbh, rofl, wtf, bff, wyd, lylc, brb, atm, imao, sml, btw, bw, imho, fyi, ppl, sob, ttyl, imo, ltr, thx, kk, omg, omfg, ttys, afn, bbs, cya, ez, f2f, gtr, ic, jk, k, ly, ya, nm, np, plz, ru, so, tc, tmi, ym, ur, u, sol, fml Consider also <https://www.netlingo.com/acronyms.php>, if you want, for no-bonus completion.

Table 4: Miscellaneous feature category specifications.